

# Building Diabetes Prediction Models

```
diabetes = readRDS("clean_diabetes_data.rds")
```

## Model 1 Preparation

```
df_multilog <- diabetes |>
  select(diabetes_012, high_bp, high_chol, chol_check, bmi, smoker, stroke,
         heart_diseaseor_attack, phys_activity, fruits, veggies, hvy_alcohol_consump,
         any_healthcare, no_docbc_cost, gen_hlth, ment_hlth, phys_hlth, diff_walk,
         sex, age, education, income)
df_multilog
```

```
## # A tibble: 253,680 x 22
##   diabetes_012 high_bp high_chol chol_check  bmi smoker stroke
##   <fct>         <fct>   <fct>    <fct>    <dbl> <fct> <fct>
## 1 No Diabetes  Yes     Yes     Yes      40 Yes  No
## 2 No Diabetes  No      No      No      25 Yes  No
## 3 No Diabetes  Yes     Yes     Yes      28 No   No
## 4 No Diabetes  Yes     No      Yes      27 No   No
## 5 No Diabetes  Yes     Yes     Yes      24 No   No
## 6 No Diabetes  Yes     Yes     Yes      25 Yes  No
## 7 No Diabetes  Yes     No      Yes      30 Yes  No
## 8 No Diabetes  Yes     Yes     Yes      25 Yes  No
## 9 Diabetes     Yes     Yes     Yes      30 Yes  No
## 10 No Diabetes No      No      Yes      24 No   No
## # i 253,670 more rows
## # i 15 more variables: heart_diseaseor_attack <fct>, phys_activity <fct>,
## #   fruits <fct>, veggies <fct>, hvy_alcohol_consump <fct>,
## #   any_healthcare <fct>, no_docbc_cost <fct>, gen_hlth <ord>, ment_hlth <dbl>,
## #   phys_hlth <dbl>, diff_walk <fct>, sex <fct>, age <ord>, education <ord>,
## #   income <ord>
```

## train/test split

```
# using createDataPartition because diabetes_012 is unbalanced based on EDA.
train_test_split <- createDataPartition(df_multilog$diabetes_012, p = 0.8, list = FALSE)
train <- df_multilog[train_test_split, ]
test <- df_multilog[-train_test_split, ]

y_test <- test$diabetes_012
```

## F1 function for multiclass

```
multi_f1 <- function(true, pred) {
  # Convert to factors
```

```

true <- as.factor(true)
pred <- as.factor(pred)

# make sure all classes are taken into consideration (bc diabetes=1 is rare)
class <- union(levels(true), levels(pred))
true <- factor(true, levels = class)
pred <- factor(pred, levels = class)

# A vector to store F1 for each class
multi_f1_vec <- numeric(length(class))

# Loop over each class
for (j in seq_along(class)) {
  i <- class[j]
  tp <- sum(true == i & pred == i)
  fn <- sum(true == i & pred != i)
  fp <- sum(true != i & pred == i)
  rec <- if((tp + fn) == 0) 0 else tp / (tp + fn)
  prec <- if((tp + fp) == 0) 0 else tp / (tp + fp)

  if(prec + rec == 0) {
    multi_f1_vec[j] <- 0
  } else {
    multi_f1_vec[j] <- 2 * prec * rec / (prec + rec)
  }
}
mean(multi_f1_vec)
}

```

## Model 1: Multinomial Logistic Regression

```

model1 <- multinom(diabetes_012 ~ ., data = train, trace = FALSE)

model1_pred <- predict(model1, newdata = test, type = "class")
model1_prob <- predict(model1, newdata = test, type = "prob")

model1_accuracy <- mean(model1_pred == y_test)
cat("Accuracy is:", model1_accuracy, "\n")

## Accuracy is: 0.8481719

model1_macro_f1 <- multi_f1(y_test, model1_pred)
cat("Macro F1 score is:", model1_macro_f1, "\n")

## Macro F1 score is: 0.39549

model1_macro_auc <- as.numeric(multiclass.roc(response = y_test, predictor = model1_prob)$auc)
cat("Macro AUC is:", model1_macro_auc, "\n")

## Macro AUC is: 0.6955571

```

### 1.1 ROC

```

cols <- c("blue", "green", "pink")

```

```

# Compute one-vs-rest ROC and AUC for each class
modell1_roc1 <- roc(as.numeric(y_test == levels(y_test)[1]), modell1_prob[, levels(y_test)[1]])

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
modell1_roc2 <- roc(as.numeric(y_test == levels(y_test)[2]), modell1_prob[, levels(y_test)[2]])

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
modell1_roc3 <- roc(as.numeric(y_test == levels(y_test)[3]), modell1_prob[, levels(y_test)[3]])

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

# Plot ROC
plot(modell1_roc1, col = cols[1], main = "ROC - Multinomial Logistic Regression")
lines(modell1_roc2, col = cols[2])
lines(modell1_roc3, col = cols[3])

# Compute class AUCs
modell1_class_aucs <- sapply(levels(y_test), function(cl) round(auc(roc(as.numeric(y_test == cl), modell1_


## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

# Add legend with class AUCs + macro AUC
legend("bottomright",
      legend = c(paste0("Class ", levels(y_test), " (AUC=", modell1_class_aucs, ")"),
                paste0("Macro AUC=", round(modell1_macro_auc, 3)),
                paste0("Accuracy=", round(modell1_accuracy, 3)),
                paste0("Macro F1 Score=", round(modell1_macro_f1, 3))),
      col = c(cols, rep(NA, 3)),
      lwd = c(rep(2, 3), rep(NA, 3)))

```



## 1.2 Confusion matrix heatmap


```

model1_cm <- table(True = y_test, Pred = model1_pred)
model1_cm_df <- as.data.frame(model1_cm) |>
  rename(True = True, Predicted = Pred, Count = Freq)

ggplot(model1_cm_df, aes(x = Predicted, y = True, fill = Count)) +
  geom_tile(color = "white") +
  geom_text(aes(label = Count), color = "black", size = 5) +
  scale_fill_gradient(low = "white", high = "purple") +
  labs(title = "Confusion Matrix Heatmap - Multi Log Regression", x = "Predicted", y = "True") +
  theme_minimal(base_size = 14)

```

## Confusion Matrix Heatmap – Multi Log Regressior



## Model 2 Preparation

```
df_ranger <- diabetes |>
  select(diabetes_2, high_bp, high_chol, chol_check, bmi, smoker, stroke,
         heart_diseaseor_attack, phys_activity, fruits, veggies, hvy_alcohol_consump,
         any_healthcare, no_docbc_cost, gen_hlth, ment_hlth, phys_hlth, diff_walk,
         sex, age, education, income)

df_ranger$diabetes_2 <- factor(df_ranger$diabetes_2,
                              levels = c(0, 1),
                              labels = c("No_Diabetes", "Diabetes"))

df_ranger
```

```
## # A tibble: 253,680 x 22
##   diabetes_2 high_bp high_chol chol_check  bmi smoker stroke
##   <fct>      <fct>   <fct>   <fct>    <dbl> <fct> <fct>
## 1 No_Diabetes Yes     Yes     Yes      40 Yes  No
## 2 No_Diabetes No      No      No       25 Yes  No
## 3 No_Diabetes Yes     Yes     Yes      28 No   No
## 4 No_Diabetes Yes     No      Yes      27 No   No
## 5 No_Diabetes Yes     Yes     Yes      24 No   No
## 6 No_Diabetes Yes     Yes     Yes      25 Yes  No
## 7 No_Diabetes Yes     No      Yes      30 Yes  No
## 8 No_Diabetes Yes     Yes     Yes      25 Yes  No
## 9 Diabetes   Yes     Yes     Yes      30 Yes  No
## 10 No_Diabetes No      No      Yes      24 No   No
## # i 253,670 more rows
```

```
## # i 15 more variables: heart_diseaseor_attack <fct>, phys_activity <fct>,
## #   fruits <fct>, veggies <fct>, hvy_alcohol_consump <fct>,
## #   any_healthcare <fct>, no_docbc_cost <fct>, gen_hlth <ord>, ment_hlth <dbl>,
## #   phys_hlth <dbl>, diff_walk <fct>, sex <fct>, age <ord>, education <ord>,
## #   income <ord>
```

## Model 2: Binary Ranger/RandomForest Model

### 2.1 Data splitting

```
set.seed(1)

index <- createDataPartition(df_ranger$diabetes_2, p = 0.8, list = FALSE)
train_data <- df_ranger[index, ]
test_data <- df_ranger[-index, ]

y_test <- test_data$diabetes_2
```

### 2.2 Train Binary Random Forest Model (using ranger)

```
rf_model <- ranger(
  dependent.variable.name = "diabetes_2",
  data = train_data,
  num.trees = 500,
  importance = "impurity",
  seed = 1,
  verbose = TRUE,
  probability = TRUE
)
```

```
## Growing trees.. Progress: 64%. Estimated remaining time: 17 seconds.
```

### 2.3 Prediction and Metrics

```
pred <- predict(rf_model, data = test_data, type = "response")

# Probability of Diabetes (positive class)
model_prob <- pred$predictions[, "Diabetes"]

# Class predictions = class with highest probability
model_pred <- colnames(pred$predictions)[ max.col(pred$predictions) ]
model_pred <- factor(model_pred, levels = levels(y_test))

# Confusion matrix
cm <- confusionMatrix(model_pred, y_test)

# Accuracy
model_accuracy <- as.numeric(cm$overall['Accuracy'])

# F1 for Diabetes
f1_diabetes <- as.numeric(cm$byClass['F1'])
# F1 for No_Diabetes
cm_rev <- confusionMatrix(model_pred, y_test, positive = "No_Diabetes")
f1_no_diabetes <- as.numeric(cm_rev$byClass['F1'])
```

```

# Macro F1
model_macro_f1 <- mean(c(f1_diabetes, f1_no_diabetes))

# AUC for Diabetes
model_roc <- roc(response = y_test, predictor = model_prob)

## Setting levels: control = No_Diabetes, case = Diabetes
## Setting direction: controls < cases
model_auc <- as.numeric(model_roc$auc)

cat("Accuracy is:      ", round(model_accuracy, 4), "\n")

## Accuracy is:      0.8523
cat("Macro F1 score is:", round(model_macro_f1, 4), "\n")

## Macro F1 score is: 0.9177
cat("ROC AUC (Binary) is:", round(model_auc, 4), "\n")

## ROC AUC (Binary) is: 0.8195

```

## 2.4 Confusion matrix heatmap

```

model_cm_table <- cm$table


model_cm_df <- as.data.frame(model_cm_table) |>
  rename(True = Reference, Predicted = Prediction, Count = Freq)

p_heatmap <- ggplot(model_cm_df, aes(x = Predicted, y = True, fill = Count)) +
  geom_tile(color = "white") +
  geom_text(aes(label = Count), color = "black", size = 5) +
  scale_fill_gradient(low = "#e0f2f1", high = "#00796b") +
  labs(
    title = "Confusion Matrix Heatmap - Random Forest",
    x = "Predicted Class",
    y = "True Class"
  ) +
  theme_minimal(base_size = 14) +
  theme(plot.title = element_text(hjust = 0.5))

print(p_heatmap)

```

## Confusion Matrix Heatmap – Random Forest




### 2.5 ROC curve

```
plot(model_roc,
     col = "darkred",
     main = "ROC Curve - Binary Random Forest",
     lwd = 2)

abline(a = 0, b = 1, lty = 2, col = "grey50")

legend(
  "bottomright",
  legend = c(
    paste0("Diabetes vs No Diabetes (AUC=", round(model_auc, 3), ")"),
    paste0("Accuracy=", round(model_accuracy, 3)),
    paste0("Macro F1 Score=", round(model_macro_f1, 3))
  ),
  col = c("darkred", rep(NA, 2)),
  lwd = c(2, rep(NA, 2))
)
```






## 2.6 Variable importance

```
importance_df <- data.frame(
  Variable = names(rf_model$variable.importance),
  Importance = rf_model$variable.importance
) |>
  arrange(desc(Importance))

ggplot(importance_df[1:15, ], aes(x = reorder(Variable, Importance), y = Importance)) +
  geom_col(fill = "#00796b") +
  coord_flip() +
  labs(title = "Top 15 Variable Importances (Random Forest)",
       x = "Variable",
       y = "Importance") +
  theme_minimal(base_size = 14)
```

## Top 15 Variable Importances (Random Forest)



### Model 3: Binary Logistic Regression Model

```
diabetes_temp = diabetes %>%
  mutate(
    diabetes_binary = ifelse(diabetes_012 == "No Diabetes", 0, 1),
    diabetes_binary = as.numeric(diabetes_binary)
  )
```

We fit a logistic regression model using health and lifestyle predictors such as blood pressure, cholesterol, BMI, smoking status, physical activity, age, and sex.

#### 3.1 Fit the Model

```
logit_model = glm(
  diabetes_binary ~ high_bp + high_chol + bmi + smoker + phys_activity + age + sex,
  data = diabetes_temp,
  family = binomial(link = "logit")
)
```

```
summary(logit_model)
```

```
##
## Call:
## glm(formula = diabetes_binary ~ high_bp + high_chol + bmi + smoker +
##     phys_activity + age + sex, family = binomial(link = "logit"),
##     data = diabetes_temp)
##
```

```
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -4.9335132  0.0337506 -146.175 < 2e-16 ***
## high_bpYes     0.9334468  0.0133878  69.724 < 2e-16 ***
## high_cholYes   0.6788247  0.0125063  54.279 < 2e-16 ***
## bmi            0.0723489  0.0008718  82.984 < 2e-16 ***
## smokerYes      0.1132249  0.0119856   9.447 < 2e-16 ***
## phys_activityYes -0.3568117  0.0128025 -27.870 < 2e-16 ***
## age.L          2.4501416  0.0602296  40.680 < 2e-16 ***
## age.Q          -0.5788621  0.0574324 -10.079 < 2e-16 ***
## age.C          -0.1757155  0.0538836  -3.261 0.00111 **
## age^4           0.0511402  0.0511210   1.000 0.31713
## age^5          -0.1167826  0.0483849  -2.414 0.01580 *
## age^6           0.0237069  0.0449595   0.527 0.59799
## age^7           0.0624394  0.0414595   1.506 0.13206
## age^8          -0.0433950  0.0375788  -1.155 0.24818
## age^9           0.0397879  0.0332257   1.198 0.23111
## age^10          -0.0698033  0.0289679  -2.410 0.01597 *
## age^11           0.0642426  0.0254442   2.525 0.01158 *
## age^12           0.0188388  0.0217445   0.866 0.38629
## sexMale         0.1436624  0.0119704  12.001 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 221031  on 253679  degrees of freedom
## Residual deviance: 185540  on 253661  degrees of freedom
## AIC: 185578
##
## Number of Fisher Scoring iterations: 6
```

### 3.2 Odds Ratios and 95% Confidence Intervals

```
exp(cbind(
  OR = coef(logit_model),
  confint.default(logit_model)
))
```

```
##               OR          2.5 %       97.5 %
## (Intercept)   0.007201159 0.006740217 0.007693625
## high_bpYes     2.543260320 2.477394210 2.610877603
## high_cholYes   1.971559144 1.923819868 2.020483062
## bmi            1.075030367 1.073194950 1.076868923
## smokerYes      1.119883711 1.093882640 1.146502816
## phys_activityYes 0.699904251 0.682560422 0.717688787
## age.L          11.589987419 10.299483382 13.042188952
## age.Q           0.560535820 0.500860543 0.627321138
## age.C           0.838856596 0.754782668 0.932295373
## age^4           1.052470431 0.952128709 1.163386837
## age^5           0.889778609 0.809276024 0.978289173
## age^6           1.023990135 0.937618435 1.118318237
## age^7           1.064429908 0.981356280 1.154535872
## age^8           0.957533083 0.889542471 1.030720438
```

```
## age^9          1.040590090  0.974984879  1.110609773
## age^10         0.932577285  0.881104300  0.987057255
## age^11         1.066351021  1.014476588  1.120878012
## age^12         1.019017415  0.976500927  1.063385056
## sexMale        1.154494301  1.127723237  1.181900884
```

### 3.3 Predicted Probability, Class Prediction, and Confusion Matrix

```
# Predicted probability
diabetes_temp$pred_prob <- predict(logit_model, type = "response")

# Convert to class prediction at threshold = 0.5
diabetes_temp$pred_class <- ifelse(diabetes_temp$pred_prob > 0.5, 1, 0)

# Confusion matrix
table(
  Actual = diabetes_temp$diabetes_binary,
  Predicted = diabetes_temp$pred_class
)
```

```
##      Predicted
## Actual      0      1
##      0 210761  2942
##      1  36471  3506
```


### 3.4 ROC Curve and AUC

```
library(pROC)

roc_obj <- roc(diabetes_temp$diabetes_binary, diabetes_temp$pred_prob)

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
auc(roc_obj)

## Area under the curve: 0.7845
plot(roc_obj, main = "ROC Curve for Logistic Regression")
```



## Model 4: Binary XGBoost

We trained an XGBoost model using a 70/30 train-test split. Predictors were converted to numeric matrices using `model.matrix`.

### 4.1 Train/Test Split and Data Preparation

```
library(xgboost)
library(caret)
library(dplyr)

#Train/Test Split
set.seed(123)

train_index = createDataPartition(diabetes_temp$diabetes_binary, p = 0.7, list = FALSE)
train=diabetes_temp[train_index, ]
test=diabetes_temp[-train_index, ]

train$binary_outcome=as.numeric(train$diabetes_binary)
test$binary_outcome=as.numeric(test$diabetes_binary)

x_train <- model.matrix(diabetes_binary ~ . - 1, data = train)
x_test  <- model.matrix(diabetes_binary ~ . - 1, data = test)

dtrain=xgb.DMatrix(data = as.matrix(x_train), label = train$diabetes_binary)
dtest=xgb.DMatrix(data = as.matrix(x_test), label = test$diabetes_binary)
```

## 4.2 Train the XGBoost Model

```
# Set Hyperparameters

params = list(
  objective = "binary:logistic",
  eval_metric = "auc",
  max_depth = 4,
  eta = 0.1,
  subsample = 0.8,
  colsample_bytree = 0.8
)

# Train Model

xgb_model = xgb.train(
  params = params,
  data = dtrain,
  nrounds = 200,
  watchlist = list(train = dtrain, test = dtest),
  verbose = 0
)
```

## 4.3 Predictions and Confusion Matrix

```
pred_prob <- predict(xgb_model, dtest)
pred_class <- ifelse(pred_prob >= 0.5, 1, 0)

confusionMatrix(
  factor(pred_class),
  factor(test$diabetes_binary),
  positive = "1"
)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##      0 64016      0
##      1      0 12088
##
##              Accuracy : 1
##              95% CI : (1, 1)
##      No Information Rate : 0.8412
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 1
##
##      McNemar's Test P-Value : NA
##
##              Sensitivity : 1.0000
##              Specificity : 1.0000
##      Pos Pred Value : 1.0000
##      Neg Pred Value : 1.0000
```

```
##           Prevalence : 0.1588
##           Detection Rate : 0.1588
##           Detection Prevalence : 0.1588
##           Balanced Accuracy : 1.0000
##
##           'Positive' Class : 1
##
```

#### 4.4 AUC

```
xgb_auc = roc(test$diabetes_binary, pred_prob)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
auc(xgb_auc)
```

```
## Area under the curve: 1
```

#### 4.5 Variable Importance Plot

```
importance_matrix <- xgb.importance(model = xgb_model)
xgb.plot.importance(importance_matrix)
```

